


GoogLeNet

Kyeong Hwan Moon

01 GoogleNet

Introduction



01 Dataset

Inspection



n02097047 (196)



n01682714 (40)



n03134739 (522)



n04254777 (806)



n02859443 (449)



n02096177 (192)



n02107683 (239)



n01443537 (1)



n02264363 (318)




Icecream - ILSVRC12 (IN)



Icecream - WIN




Icecream - WINC



ILSVR red, WINC blue




Pillow - ILSVRC12 (IN)



Pillow - WIN



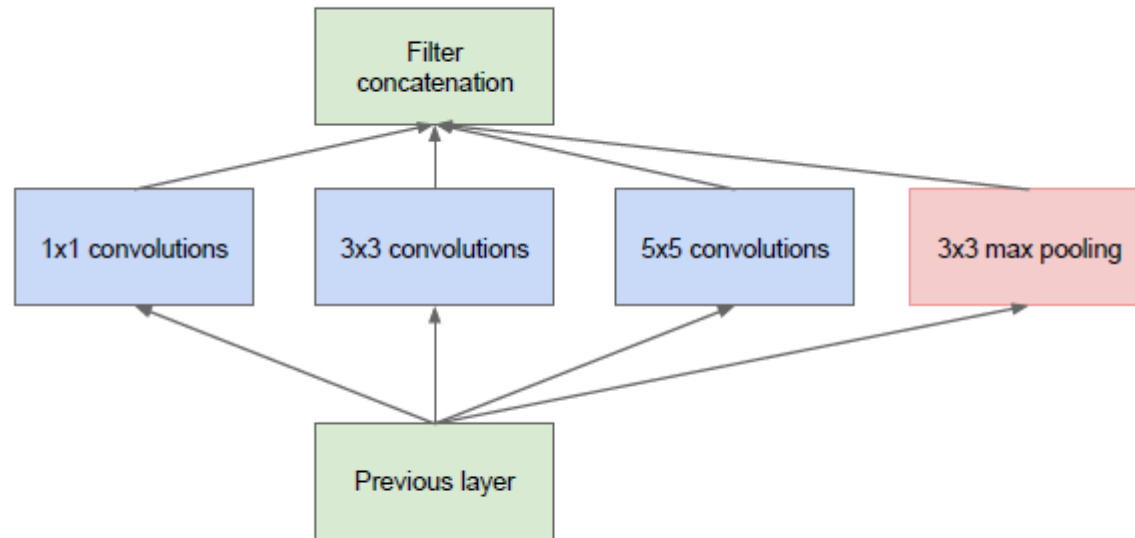
Pillow - WINC



ILSVR red, WINC blue

01 Inception Module

Naïve version





(a) Inception module, naïve version

- Increase in the number of outputs from stage to stage
- Calculation could be expensive on the top of convolutional layer

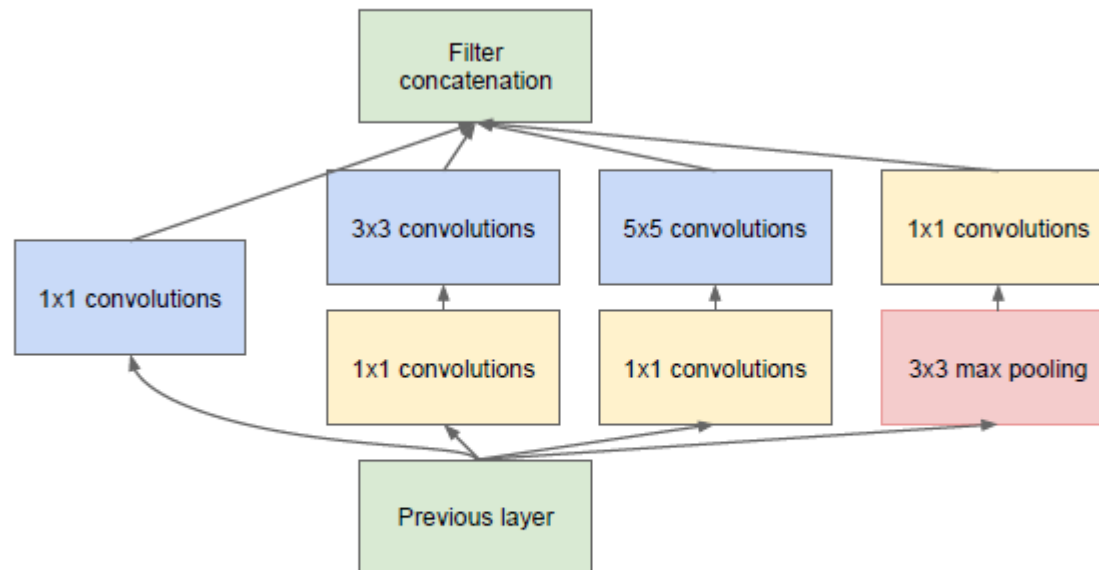
01 Inception Module

Naïve version



01 Inception Module

Dimension reduction



(b) Inception module with dimension reductions

- Used 1x1 convolutional filter
- Calculation could be cheaper than naïve version
- Used Linear activation(But ReLU could be better)

01 GoogLeNet

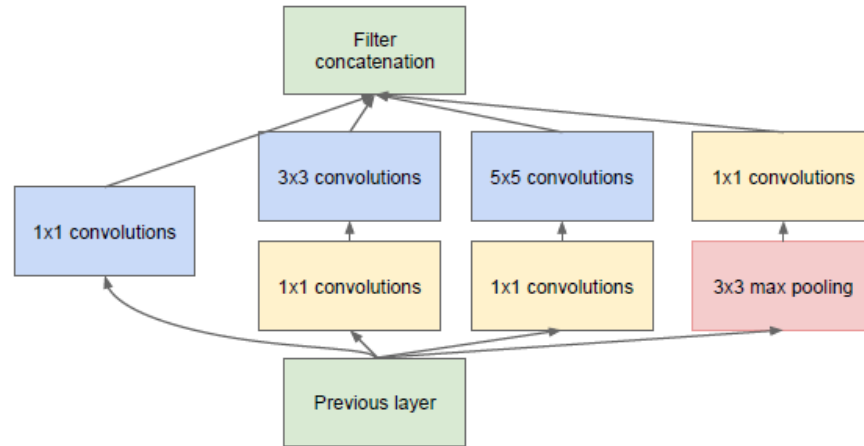
Table

type	patch size/ stride	output size	depth	$\#1 \times 1$	$\#3 \times 3$ reduce	$\#3 \times 3$	$\#5 \times 5$ reduce	$\#5 \times 5$	pool proj	params	ops
convolution	$7 \times 7 / 2$	$112 \times 112 \times 64$	1							2.7K	34M
max pool	$3 \times 3 / 2$	$56 \times 56 \times 64$	0								
convolution	$3 \times 3 / 1$	$56 \times 56 \times 192$	2		64	192				112K	360M
max pool	$3 \times 3 / 2$	$28 \times 28 \times 192$	0								
inception (3a)		$28 \times 28 \times 256$	2	64	96	128	16	32	32	159K	128M
inception (3b)		$28 \times 28 \times 480$	2	128	128	192	32	96	64	380K	304M
max pool	$3 \times 3 / 2$	$14 \times 14 \times 480$	0								
inception (4a)		$14 \times 14 \times 512$	2	192	96	208	16	48	64	364K	73M
inception (4b)		$14 \times 14 \times 512$	2	160	112	224	24	64	64	437K	88M
inception (4c)		$14 \times 14 \times 512$	2	128	128	256	24	64	64	463K	100M
inception (4d)		$14 \times 14 \times 528$	2	112	144	288	32	64	64	580K	119M
inception (4e)		$14 \times 14 \times 832$	2	256	160	320	32	128	128	840K	170M
max pool	$3 \times 3 / 2$	$7 \times 7 \times 832$	0								
inception (5a)		$7 \times 7 \times 832$	2	256	160	320	32	128	128	1072K	54M
inception (5b)		$7 \times 7 \times 1024$	2	384	192	384	48	128	128	1388K	71M
avg pool	$7 \times 7 / 1$	$1 \times 1 \times 1024$	0								
dropout (40%)		$1 \times 1 \times 1024$	0								
linear		$1 \times 1 \times 1000$	1							1000K	1M
softmax		$1 \times 1 \times 1000$	0								

Table 1: GoogLeNet incarnation of the Inception architecture


01 GoogLeNet

GoogLeNet



- Large depth of the network
- Auxiliary classifiers are calculated while training
- Their loss gets added to the total loss of the network with a discount weight

(b) Inception module with dimension reductions



01 Code

```
class Inception(nn.Module):
    def __init__(self, in_planes, kernel_1_out, kernel_3_in, kernel_3_out,
                  kernel_5_in, kernel_5_out, pool_planes):
        super(Inception, self).__init__()

        # 1x1 conv block
        self.block1 = nn.Sequential(
            nn.Conv2d(in_planes, kernel_1_out, kernel_size=1),
            nn.ReLU(inplace=True),
        )

        # 1x1 conv block -> 3x3 conv block
        self.block2 = nn.Sequential(
            nn.Conv2d(in_planes, kernel_3_in, kernel_size=1),
            nn.ReLU(inplace=True),
            nn.Conv2d(kernel_3_in, kernel_3_out, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),
        )

        # 1x1 conv block -> 5x5 conv block
        self.block3 = nn.Sequential(
            nn.Conv2d(in_planes, kernel_5_in, kernel_size=1),
            nn.ReLU(inplace=True),
            nn.Conv2d(kernel_5_in, kernel_5_out, kernel_size=5, padding=2),
            nn.ReLU(inplace=True),
        )
```

```
        # 3x3 pool block -> 1x1 conv block
        self.block4 = nn.Sequential(
            nn.MaxPool2d(kernel_size=3, stride=1, padding=1),
            nn.Conv2d(in_planes, pool_planes, kernel_size=1),
            nn.ReLU(inplace=True),
        )

    def forward(self, x):
        x1 = self.block1(x)
        x2 = self.block2(x)
        x3 = self.block3(x)
        x4 = self.block4(x)
        return torch.cat([x1, x2, x3, x4], 1)
```

01 Code

```
class BasicConv2d(torch.nn.Module):
    def __init__(self, in_channels, out_channels, **kwargs):
        super(BasicConv2d, self).__init__()
        self.conv = torch.nn.Conv2d(in_channels, out_channels, **kwargs)
        self.batchnorm = torch.nn.BatchNorm2d(out_channels)
        self.relu = torch.nn.ReLU(inplace=True)

    def forward(self, x):
        x = self.conv(x)
        x = self.batchnorm(x)
        x = self.relu(x)
        return x
```

```
class Aux(torch.nn.Module):
    def __init__(self, in_channels, num_classes):
        super(Aux, self).__init__()
        self.avgpool = torch.nn.AvgPool2d(kernel_size=2, stride=2)
        self.conv = BasicConv2d(in_channels, 128, kernel_size=1)
        self.fc1 = torch.nn.Sequential(torch.nn.Linear(2 * 2 * 128, 256))
        self.fc2 = torch.nn.Linear(256, num_classes)

    def forward(self, x):
        out = self.avgpool(x)
        out = self.conv(out)
        out = out.view(out.size(0), -1)
        out = torch.nn.functional.dropout(out, 0.5, training=self.training)
        out = torch.nn.functional.relu(self.fc1(out), inplace=True)
        out = torch.nn.functional.dropout(out, 0.5, training=self.training)
        out = self.fc2(out)
        return out
```

01 Code

```
class GoogLeNet(nn.Module):
    def __init__(self, base_dim, num_classes=2, batch_size=50, aux_logits=True):
        super(GoogLeNet, self).__init__()
        self.num_classes = num_classes
        self.batch_size = batch_size

        self.post_inception = nn.Sequential(
            nn.Conv2d(3, base_dim, kernel_size=7, padding=1),
            nn.MaxPool2d(kernel_size=3, stride=2, padding=1),
            nn.Conv2d(base_dim, base_dim*3, 3, 1, 1),
            nn.MaxPool2d(3, 2, 1),
        )

        self.inception_layer1 = nn.Sequential(
            Inception(base_dim*3, 64, 96, 128, 16, 32, 32),
            Inception(base_dim*4, 128, 128, 192, 32, 96, 64),
            nn.MaxPool2d(3, 2, 1),
            Inception(480, 192, 96, 208, 16, 48, 64),
        )

        self.inception_layer2 = nn.Sequential(
            Inception(512, 160, 112, 224, 24, 64, 64),
            Inception(512, 128, 128, 256, 24, 64, 64),
            Inception(512, 112, 144, 288, 32, 64, 64),
        )
```

```
        if self.aux_logits:
            self.aux1 = Aux(512, num_classes)
            self.aux2 = Aux(528, num_classes)

        self.avgpool = torch.nn.AdaptiveAvgPool2d((1, 1))
        self.dropout = torch.nn.Dropout(0.4)
        self.fc = torch.nn.Linear(1024, num_classes)

    def forward(self, x):
        out = self.post_inception(x)
        out = self.inception_layer1(out)

        if self.training and self.aux_logits:
            aux1 = self.aux1(out)

        out = self.inception_layer2(out)

        if self.training and self.aux_logits:
            aux2 = self.aux2(out)

        out = self.inception_layer3(out)
        out = self.dropout(out)
        out = out.view(out.size(0), -1)
        out = self.fc(out)
        return out, aux1, aux2
```

Thank
You
